

# Metagenomic Classification: a Deep Learning Approach

Daniele Bellani

01 • 02 • 2018

## Some facts...

- Great excitement about **deep learning** models, due to successful applications in computer vision, speech recognition, natural language processing...
- In particular, CNN and RNN (in various flavours) obtained great results when treating **sequential data**
- **DNA** and **RNA** streams are a kind of sequential data
- First attempts on biomedical tasks, mainly involving **genomic data**
- Their application on **metagenomic data** is yet to come...

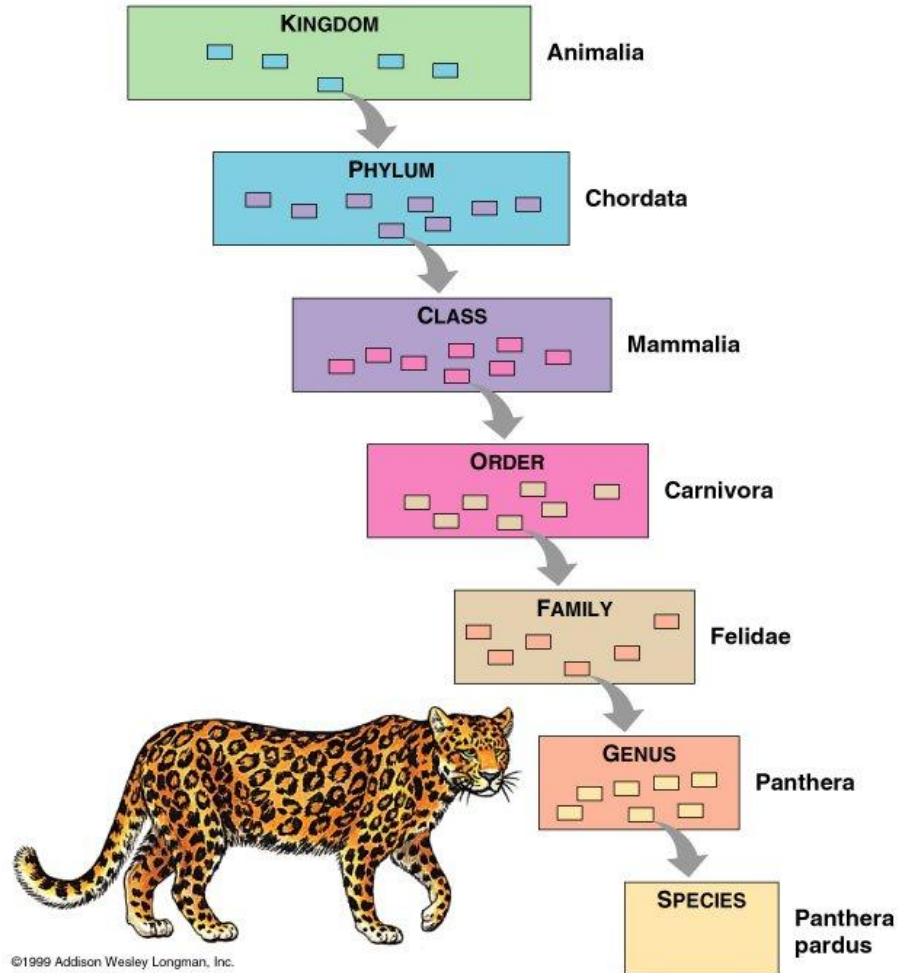
# Biological background

## Some new words...


- **Metagenomics** - study of genetic material recovered directly from **environmental** samples.
- **Reads** - fragments of genetic material.
- **Sequencing** - the process of extracting **reads** from biological samples.
- **K-mer** - DNA/RNA string of length **k**.
- **Prokaryotic cell** - type of cell which differs from the **eukaryotic cell** for not having a nucleus, having simpler internal structure and for **not** assembly in multicellular organisms (prokaryotes are **unicellular organisms**).

## Some new words...

- **Virus** - protein shell containing genetic material.
- **Microbiome** - could indicate either a population of microorganisms or the collection of their genetic material (**genomes**).
- **Taxon** - a population, or group of populations of organisms which are usually inferred to be **phylogenetically** related and share characteristics which differentiate them from other groups. Taxons are organized in a **taxonomical ranking**.




# The central(ish) dogma of genomics



# Metagenomic Classification



# An ordinary metagenomic experiment



Sampling



Sequencing



AYATCCG  
CCGGBTG  
AAWTCCT  
...



ATATCCGTCC

=

?

Classification

# Metagenomic classification

Assign a **taxon** to a **read**:

ATCCACATATTCTTTCTAATCTCATTTTTATCTACATAAAGTAAAAGTTATTCACAAAAACGTAGCTTTA



**Kingdom:** *Bacteria*


...

**Genus:** *Pelosinus*

**Specie:** *Pelosinus Fermentans*

# Why is it important? - Data Analysis


- Classification of reads is important to infer the composition of **microbial communities (microbiomes)** of the sample, and thus of the environment it comes from.
- Typical analyses relying on such operation include **pollution** analysis, **pathogens** detection, **air/water quality** analysis, etc.
- Samples can come even from the human body (we host a number of microorganism 3 times larger than the number of human cells!). These are used in **medicine** and **nutrition**.




# How to classify - Alignment

- Most of state-of-the-art procedures are based on **sequence alignment**

## Local Alignment



## Global Alignment



# How to classify - Alignment

- Each sequence produced by the sequencer serves as **query sequence**, while the **target sequence** is every sequence stored in the **genomes database**.
- The comparison can be performed on **nucleotide level** or **protein level**.



# How to classify - Alignment

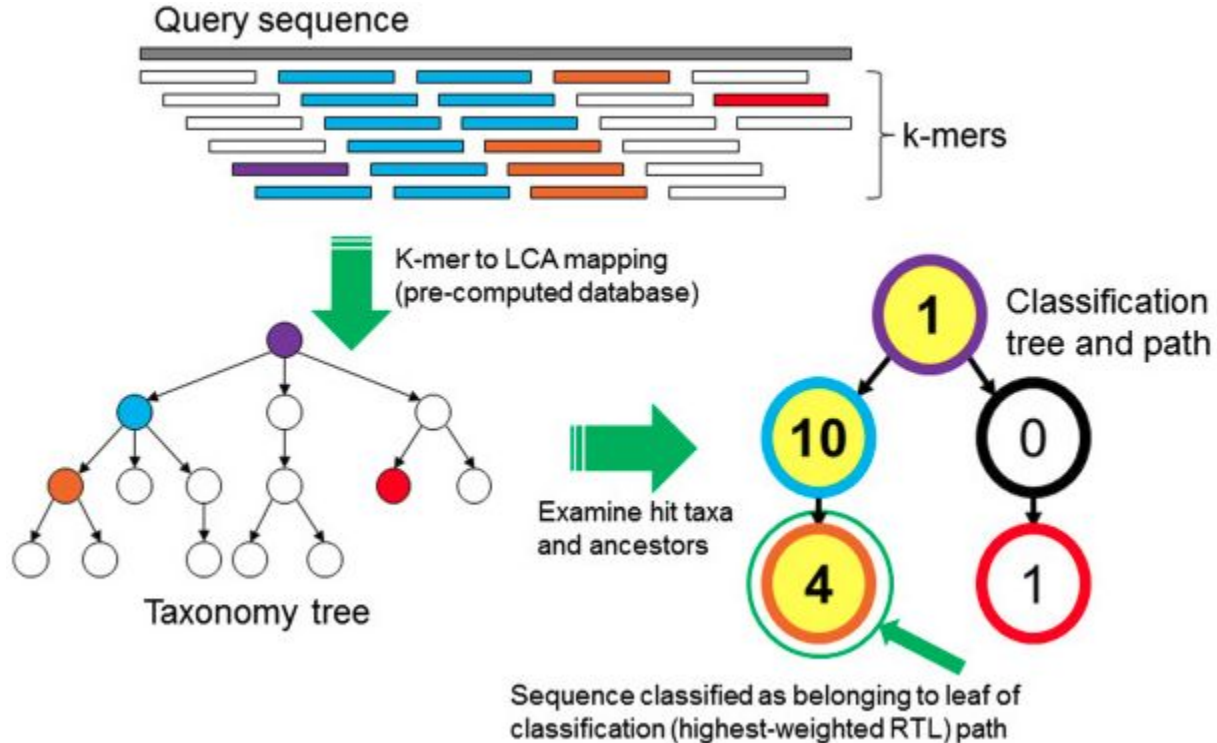
- Original dynamic programming algorithms could perform **global alignment** and **local alignment** with complexity  $O(mn)$
- Alignment algorithms have been improved in the last 30 years, achieving a linear complexity -  $O(n)$
- But alignment still remains a **time-consuming** operation:
  - Latest **Next Generation Sequencing** processes (NGS) produce billion of sequences
  - Databases contains thousands of fully-decoded genomes



# How to classify - Alignment-free methods

- In recent years, a lot of approaches not relying on alignment have been published
- These new methods are roughly distinguished in two classes:
  - **Marker gene approaches**
  - **K-mer based approaches**
- These methods have demonstrated good performances and high speed, but are weak on **real data** (due to variations) and suffer from **sampling bias**.
- The methods in these categories still rely on pre-constructed genomic databases or self-constructed mappings (these can be hundreds of GB in size!).

# How to classify - Alignment-free methods





# How to classify - Machine learning methods

- Some machine learning attempts to classification, mainly with:
  - Naive Bayes
  - SVM
- Minor attempts with:
  - Nearest Neighbor
  - Random Forest
- Machine Learning models were celebrated by biologists for their high **sensitivity (recall)**.
- Nevertheless, they have never catch the heart of biologists, due to small improvement on runtime.

A word cloud of bioinformatics tools. The word 'BLAST' is the largest and most central. Other tools are arranged around it in various sizes and colors. The colors used are red, orange, teal, and grey. The tools include: Kaiju, NBC, MG-RAST, MetaPhlAn, Kraken, LMAT, MEGAN, DectICO, Centrifuge, CLARK, mOTU, RAPSearch, GOTTCHA, CARMA3, DIAMOND, k-SLAM, Phymm, and PhymmBL.

Kaiju

NBC

MG-RAST

MetaPhlAn

Kraken

LMAT

MEGAN

DectICO

BLAST

Centrifuge

CLARK

mOTU

RAPSearch

GOTTCHA

CARMA3

DIAMOND

k-SLAM

Phymm

PhymmBL

# How to classify - Validation

- Benchmarks can be performed on both **real** or **simulated** data.
- The principal metrics for the validation of a metagenomic classifier include:
  - Accuracy
  - Precision
  - Recall (**sensitivity**)
  - Measures combining Precision and Recall (**F1, ROC**)
  - Speed (rpm)
  - Predicted vs Real microbial distributions correlation
  - Fraction of unclassified sequences (main issue on real data)
- It is important to assess the robustness of the method by executing it on different-sized reads.



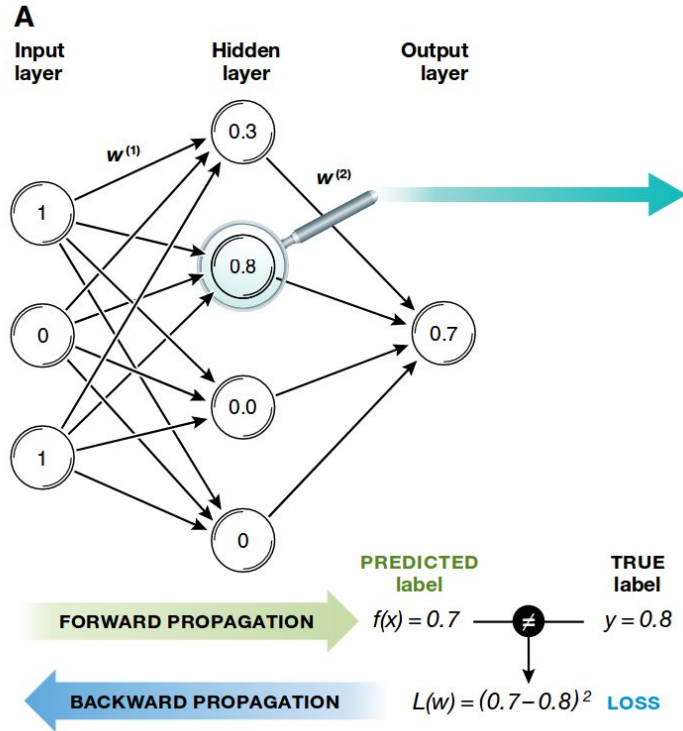
# Metagenomic classification - Recap

- Sequence classification is a fundamental part of the metagenomic pipeline
- Alignment algorithms are still the most used approach but:
  - Time-consuming
  - Dependence on databases (which may not be complete)
- Alignment-free algorithms are faster and obtain even better results but:
  - Require massive memory and disk space
  - Still dependant on genomic databases (**sampling bias**)
  - Weak on real data

**The goal is to find a method to accurately classify the highest percentage of different-sized reads from real and simulated datasets in the least time-consuming and memory-consuming way.**

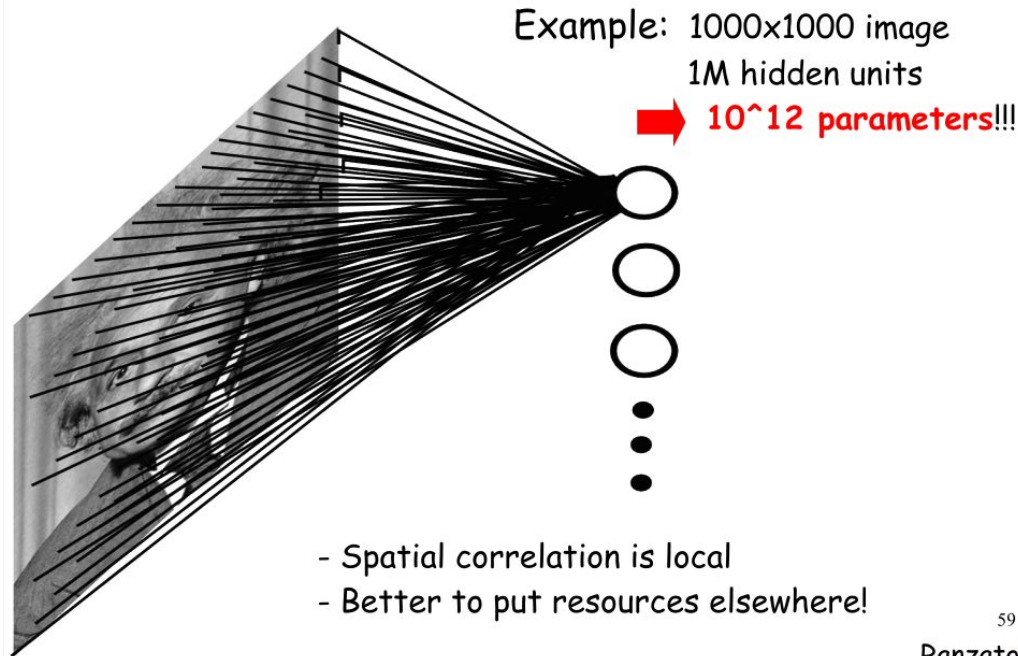
# Deep Learning for Sequential Data

# Perceptron and Multi-Layer Perceptron



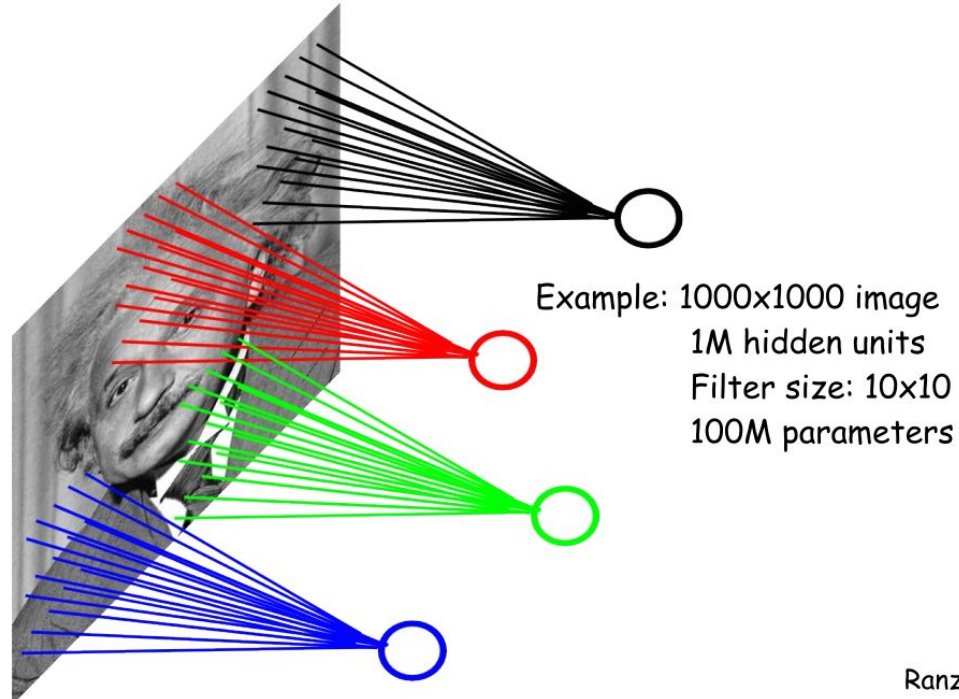
# Convolutional Neural Networks

## FULLY CONNECTED NEURAL NET



# Convolutional Neural Networks

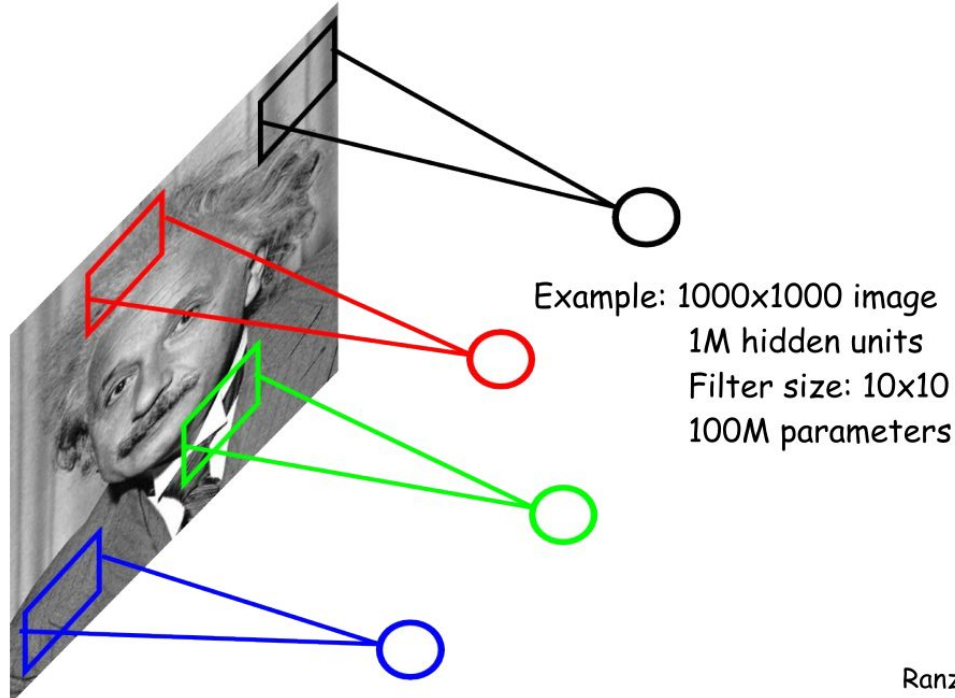
## LOCALLY CONNECTED NEURAL NET






# Convolutional Neural Networks

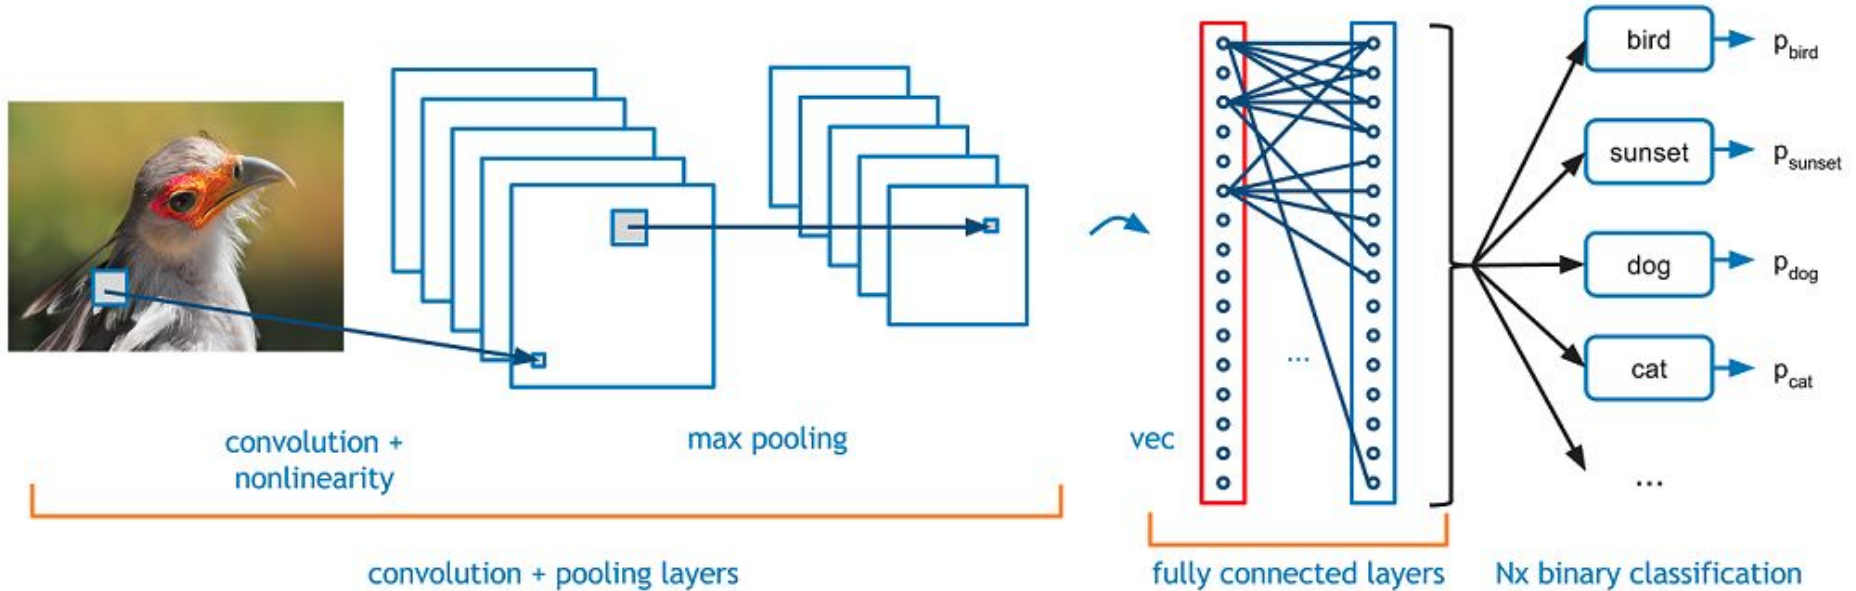
## LOCALLY CONNECTED NEURAL NET



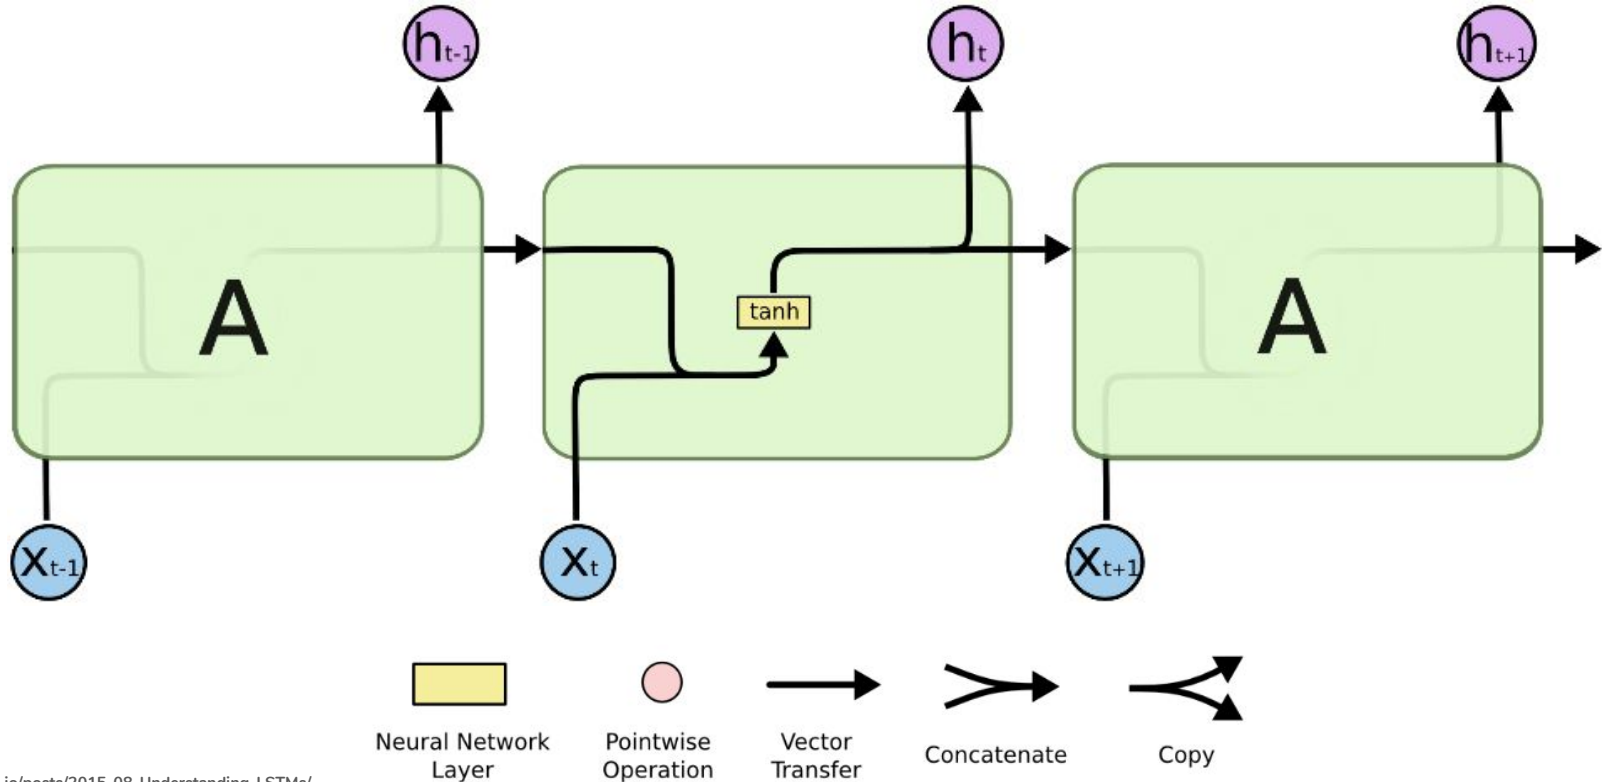
# Convolutional Neural Networks



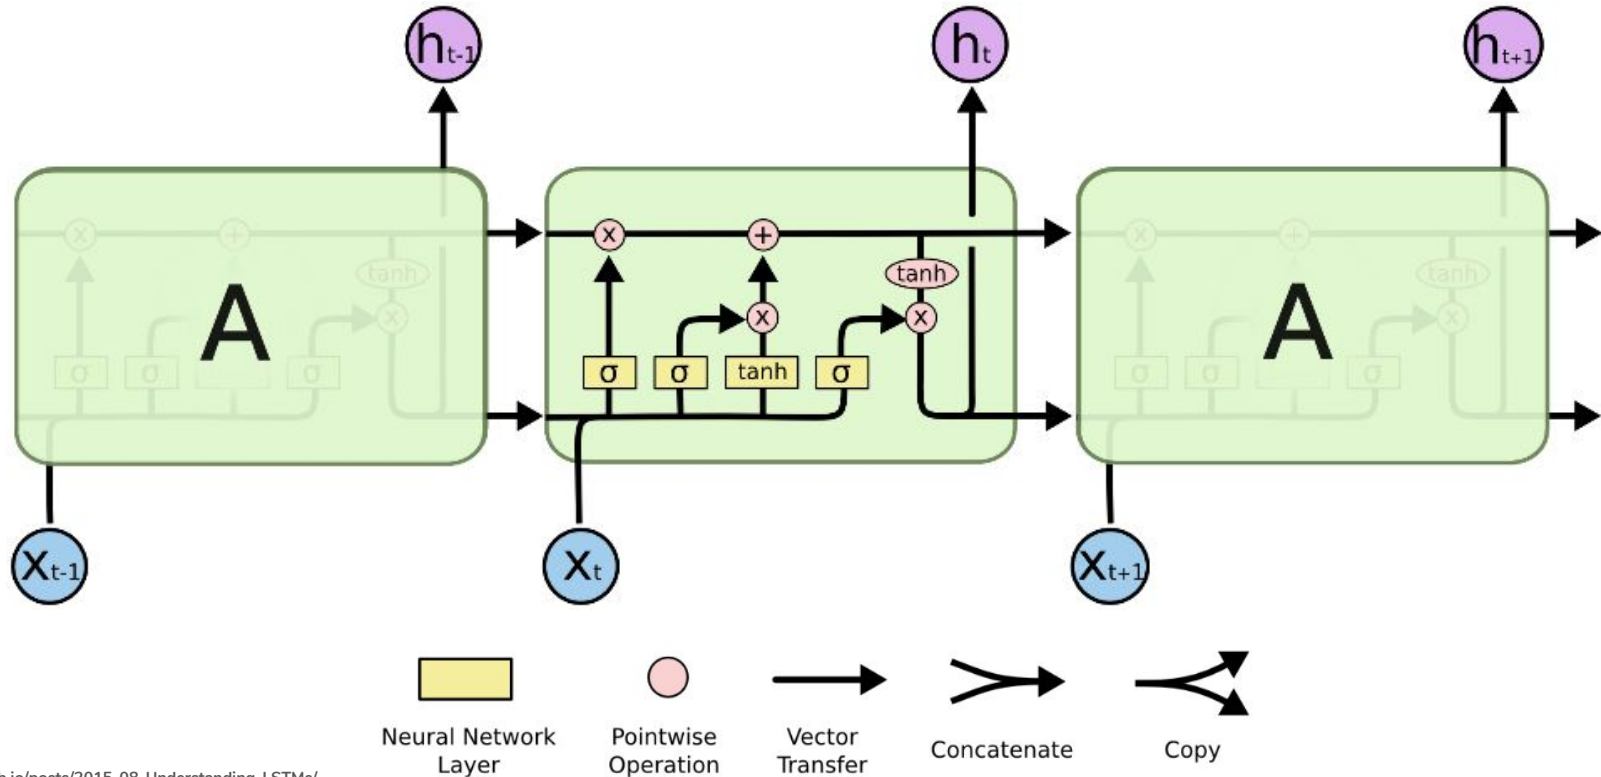
# Convolutional Neural Networks



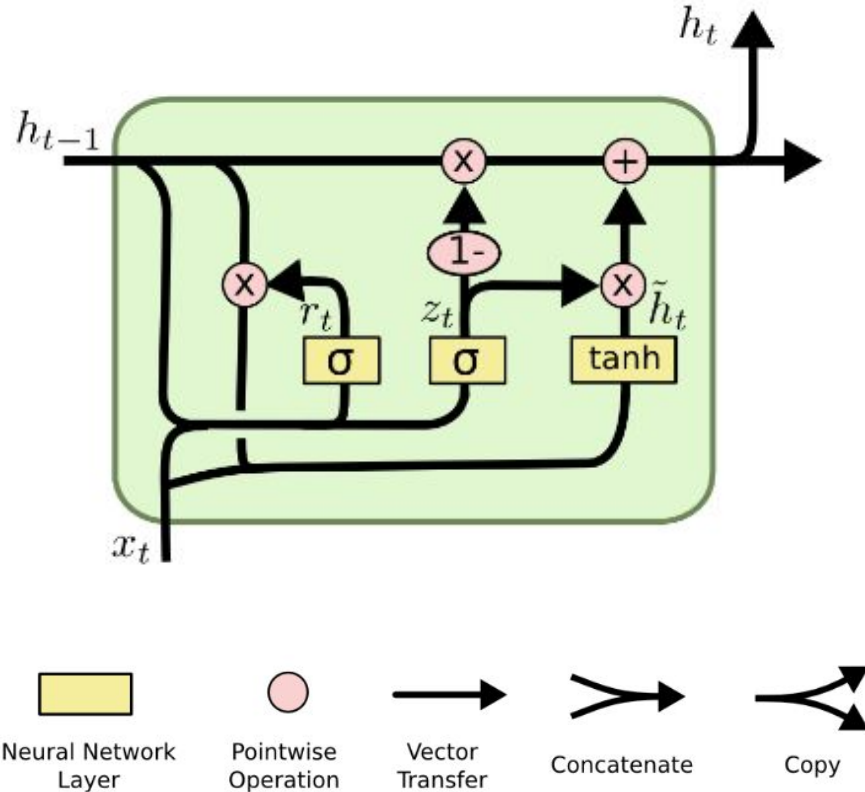
# Recurrent Neural Networks





# Recurrent Neural Networks - LSTM



# Recurrent Neural Networks - GRU



## DNN tools



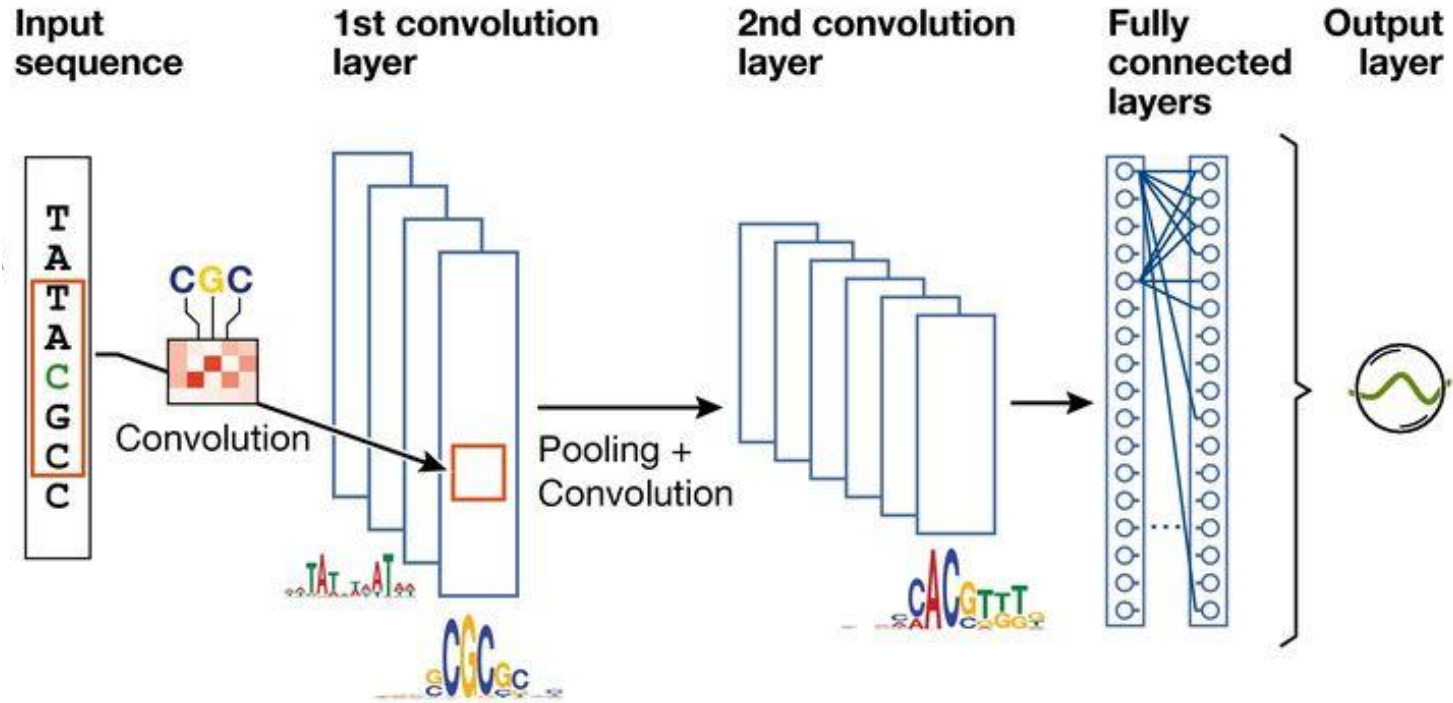
# DNN for Metagenomic Classification



# Overview

- Deep learning models have shown good performances on **image classification** and **speech recognition**
- They manage to memorize patterns hidden in input data, which can be quite complex
- Unlike k-mer based approaches, neural networks (stateful RNN above all) can treat the input sequence as a whole stream of data, making them more robust to local variations
- Unlike alignment based approaches, no comparison between sequence is needed

# Overview



# Motivations

- **Speed**
  - Heavy training phase, but prediction phase grows linearly with the input
  - No comparison between sequences
  - GPU speed-up
- **No feature extraction**
  - Both CNN and RNN take raw (yet vectorized) sequences as input
- **Genome DB independence**
  - No need for complete reconstructed genome for each taxon, but a labeled set of sequences.
- **Local variations-tolerant**
  - All the sequence points are examined and weighed, no perfect match is sought

# Motivations

- **2-level support**
  - Could work with either protein-level or nucleotide-level sequences
- **Motif discovery**
  - Possibility to explain predictions by highlighting the most significant portions of input data



# Issues

- **Class number**
  - Performance of ML algorithms decreases with the increasing number of classes
- **Unknown species**
  - Metagenomic samples are quite often full of **uncategorized** or **unexpected** species
- **Motif discovery**
  - Novel approach, still a few works, could be the toughest part
- **Parameter tuning, structure choice and learning**
  - NN are full of parameters to be optimized

# Looking around...

## Predicting effects of noncoding variants with deep learning–based sequence model

Jian Zhou<sup>1,2</sup> and Olga G Troyanskaya<sup>1,3,4</sup>


## Predicting the sequence specificities of DNA- and RNA-binding proteins by deep learning

Babak Alipanahi<sup>1,2,6</sup>, Andrew Delong<sup>1,6</sup>, Matthew T Weirauch<sup>3–5</sup> & Brendan J Frey<sup>1–3</sup>

## Convolutional neural network architectures for predicting DNA–protein binding



Haoyang Zeng, Matthew D. Edwards, Ge Liu and David K. Gifford\*

Computer Science and Artificial Intelligence Laboratory, Massachusetts Institute of Technology, Cambridge, MA 02142, USA



# SCELSE

Singapore Centre for Environmental Life Sciences Engineering



**NANYANG  
TECHNOLOGICAL  
UNIVERSITY**  

---

**SINGAPORE**

**DISC**  
DIPARTIMENTO  
DI INFORMATICA  
SISTEMISTICA  
E COMUNICAZIONE

UNIVERSITA' DEGLI STUDI  
DI MILANO  
**BICOCCA**

# Bibliography

- Breitwieser, F. P., Lu, J., & Salzberg, S. L. (2017). A review of methods and databases for metagenomic classification and assembly. Briefings in Bioinformatics.
- Menzel, P., Ng, K. L., & Krogh, A. (2016). Fast and sensitive taxonomic classification for metagenomics with Kaiju. Nature communications, 7.
- Lindgreen, S., Adair, K. L., & Gardner, P. P. (2016). An evaluation of the accuracy and speed of metagenome analysis tools. Scientific reports, 6.
- Břinda, K., Sykulski, M., & Kuchеров, G. (2015). Spaced seeds improve k-mer-based metagenomic classification. Bioinformatics, 31(22), 3584-3592.
- Ding, X., Cheng, F., Cao, C., & Sun, X. (2015). DectlCO: an alignment-free supervised metagenomic classification method based on feature extraction and dynamic selection. BMC bioinformatics, 16(1), 323.
- Angermueller, C., Pärnamaa, T., Parts, L., & Stegle, O. (2016). Deep learning for computational biology. Molecular systems biology, 12(7), 878.



# Bibliography

- Ounit, R., Wanamaker, S., Close, T. J., & Lonardi, S. (2015). CLARK: fast and accurate classification of metagenomic and genomic sequences using discriminative k-mers. *BMC genomics*, 16(1), 236.
- Marc'Aurelio Ranzato (2014), Deep Learning Tutorial  
<https://docs.google.com/viewer?a=v&pid=sites&srcid=ZGVmYXVsdGRvbWFpbnxsc3ZydHV0b3JpYWxjdnByMTR8Z3g6Njg5MmZkZTM1MDhhZWNmZA>
- Wood, D. E., & Salzberg, S. L. (2014). Kraken: ultrafast metagenomic sequence classification using exact alignments. *Genome biology*, 15(3), R46.
- Mande, S. S., Mohammed, M. H., & Ghosh, T. S. (2012). Classification of metagenomic sequences: methods and challenges. *Briefings in bioinformatics*, 13(6), 669-681.
- Rosen, G., Garbarine, E., Caseiro, D., Polikar, R., & Sokhansanj, B. (2008). Metagenome Fragment Classification Using *N*-Mer Frequency Profiles. *Advances in bioinformatics*, 2008.
- Soueidan, H., & Nikolski, M. (2017). Machine learning for metagenomics: methods and tools. *Metagenomics*, 1(1).

# Bibliography

- Quang, D., & Xie, X. (2016). DanQ: a hybrid convolutional and recurrent deep neural network for quantifying the function of DNA sequences. *Nucleic acids research*, 44(11), e107-e107.
- Lanchantin, J., Singh, R., Lin, Z., & Qi, Y. (2016). Deep motif: Visualizing genomic sequence classifications. *arXiv preprint arXiv:1605.01133*.
- Leung, M. K. K., DeLong, A., & Frey, B. J. (2017). Inference Of The Human Polyadenylation Code. *bioRxiv*, 130591.
- <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

**THANKS!**